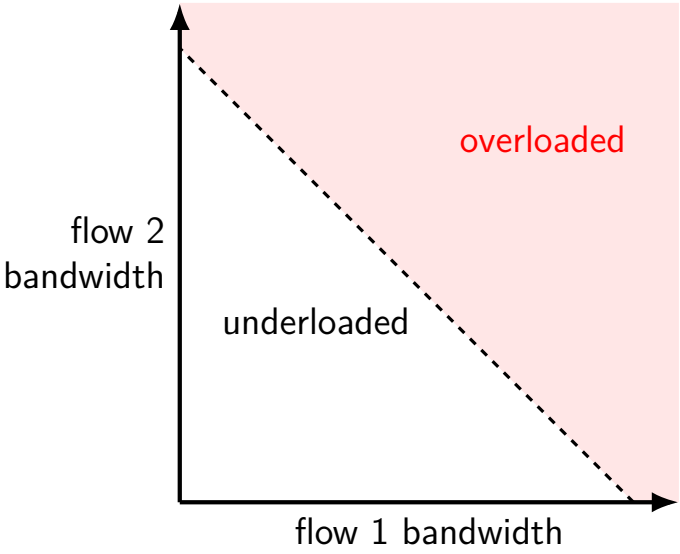
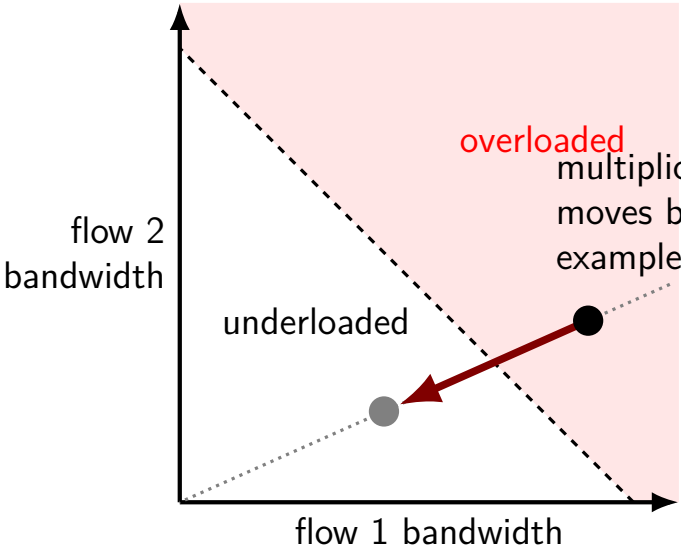
picturing sharing

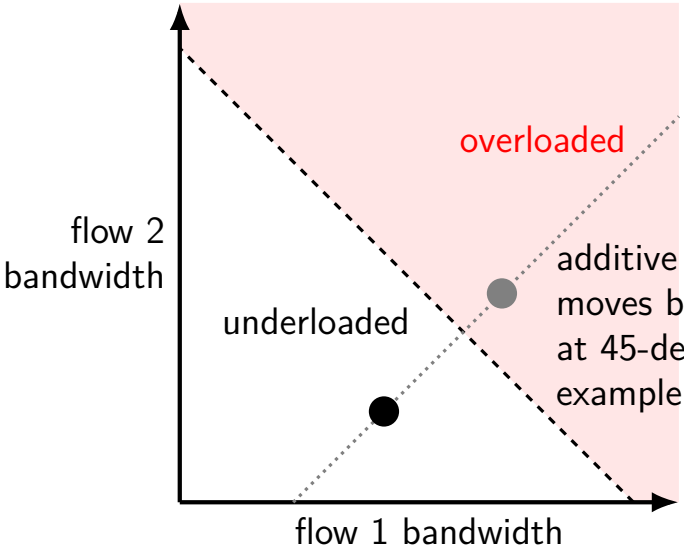


picturing sharing

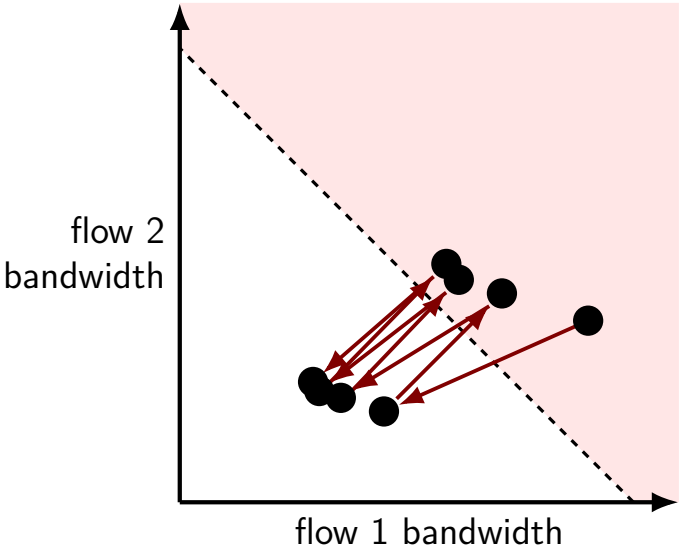


multiplicative decrease
moves bandwidths along line to origin
example: (90%, 40%) to (45%, 20%)

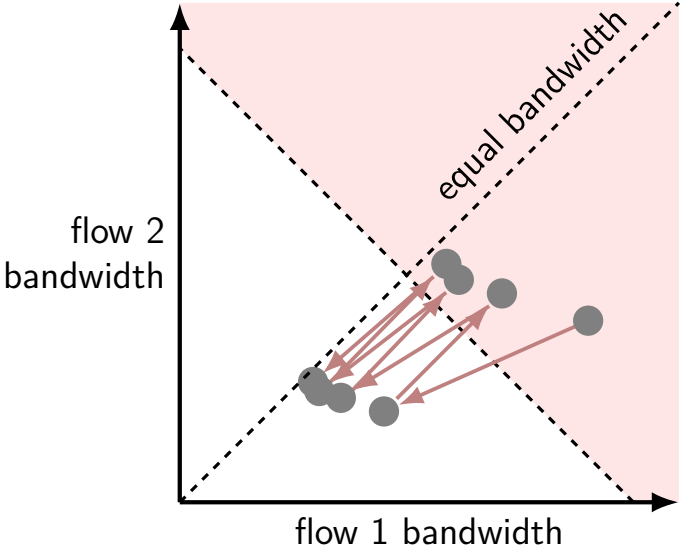
picturing sharing



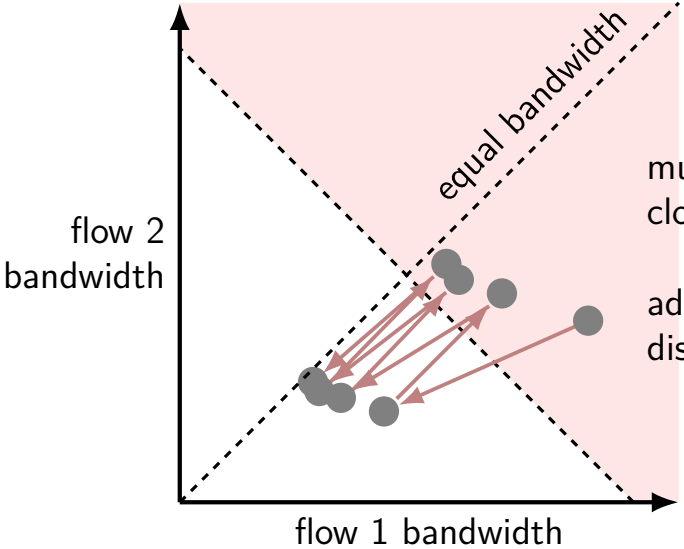
picturing sharing



picturing sharing



picturing sharing



multiplicative decrease brings closer to equal bandwidth line

additive increase keeps same distance from line

some assumptions we made

...that are probably not true

both flows experience drops equally when network overloaded

same additive increase factor (for 45 degree angle)

models that give numbers?

deciding on congestion control

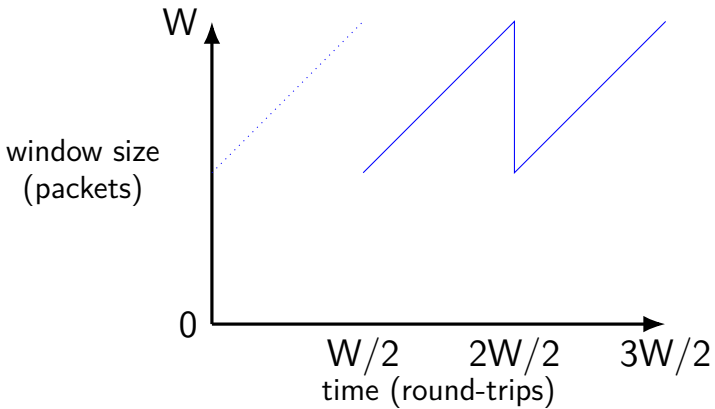
would like to do math to say how well they'll do

common approach: but gets complicated

simple example: estimating average rate of TCP-like AIMD with no congestion

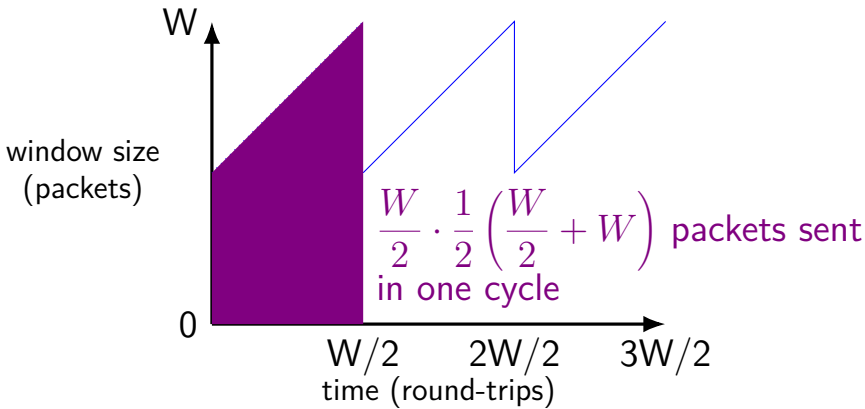
models that give numbers? (1)

figure adapted from Mathis et al, "The Macroscopic Behavior of the TCP Congestion Avoidance Algorithm"



models that give numbers? (1)

figure adapted from Mathis et al, "The Macroscopic Behavior of the TCP Congestion Avoidance Algorithm"



models that give numbers? (2)

$\frac{W}{2} \cdot \frac{1}{2}\left(\frac{W}{2} + W\right)$ per $W/2$ round trips

$\frac{3W}{4}$ packets per round trip time

packet loss should occur when window size $=$ bandwidth-delay product

at the capacity of the link(s)

so $W =$ link BW $\cdot$ RTT

$\implies$ 3/4 link BW achieved total

exercise: what things did this model miss?

some answers

RTT/delay depends on how many packets queued

packet loss could occur for other reasons

- competing connections
- network errors

'bursty' connection could trigger packet loss earlier

extra packets being sent for retransmissions

packet loss could trigger timeout/multiple decreases

behavior of other connections sharing links

delays in sending ACKs depending how fast receiver's CPU is

...

more sophisticated models?

we can add to formulas to account for other things

this is something people do, but...

most common technique is *discrete event simulation*

interlude: loss rate → transfer rate

adapted from Mathis et al, "The Macroscopic Behavior of the TCP Congestion Avoidance Algorithm"

packet loss rate $p = 1$ per (number of packets sent in W/2 round trips)

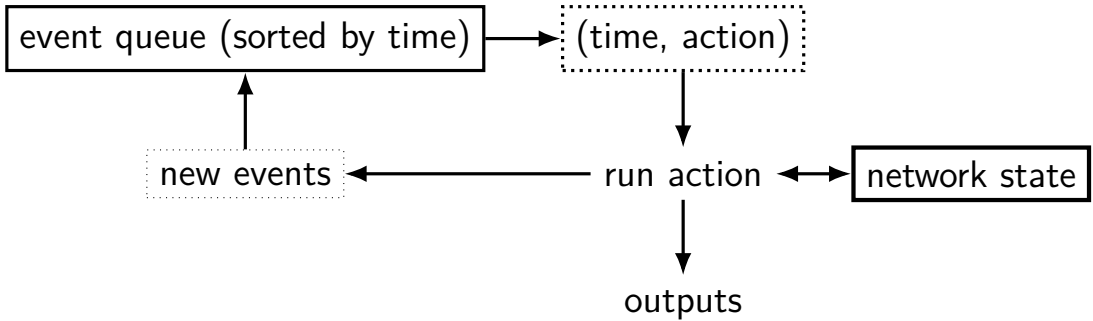
$3/4 \times W \times W/2$ packets sent in W/2 round trips

$$p = \frac{3}{8}W^2$$

solving for $W = \sqrt{\frac{8}{3p}}$

average transfer rate $= \frac{3}{4}W = C \cdot \sqrt{1/p}$ (for some C)

discrete-event simulation



action example 1

take next packet from send queue for link X

compute whether packet is lost due to error

compute when packet is done transmitting

- schedule new event to handle next packet in queue at that time

compute reception time of packet on other end of link

- schedule new event to handle packet being received at that time

action example 2

take next packet on link 0 of switch

compute next link for packet

add packet to queue for next link

schedule new events:

- to dequeue from next link (if not scheduled already)

NS-3

discrete event simulator planned for AIMD assignment

written in C++

(yes, I know it's not the most familiar language)
(obvious alternative simulators aren't in better languages...)

create simulations by writing C++ programs

an NS-3 event handler

```
bool PointToPointChannel::TransmitStart(
    Ptr<const Packet> p,
    Ptr<PointToPointNetDevice> src,
    Time txTime
) {
  // ...
  uint32_t wire = src == m_link[0].m_src ? 0 : 1;

  Simulator::ScheduleWithContext(
    m_link[wire].m_dst->GetNode()->GetId(),
    txTime + m_delay,
    &PointToPointNetDevice::Receive, m_link[wire].m_dst, p->Copy());

  // Call the tx anim callback on the net device
  m_txrxPointToPoint(p, src, ...)
  return true;
}
```

an NS-3 event handler

X::Y = Y method/variable of X class

```
bool PointToPointChannel::TransmitStart(
    Ptr<const Packet> p,
    Ptr<PointToPointNetDevice> src,
    Time txTime
) {
  // ...
  uint32_t wire = src == m_link[0].m_src ? 0 : 1;

  Simulator::ScheduleWithContext(
    m_link[wire].m_dst->GetNode()->GetId(),
    txTime + m_delay,
    &PointToPointNetDevice::Receive, m_link[wire].m_dst, p->Copy());

  // Call the tx anim callback on the net device
  m_txrxPointToPoint(p, src, ...)
  return true;
}
```

an NS-3 event handler

convention: member variables with m_
(C++ member variable ∼ Java instance variable)

```
bool PointToPointChannel::Tra
    Ptr<const Packet> p,
    Ptr<PointToPointNetDevice> src,
    Time txTime
) {
  // ...
  uint32_t wire = src == m_link[0].m_src ? 0 : 1;

  Simulator::ScheduleWithContext(
    m_link[wire].m_dst->GetNode()->GetId(),
    txTime + m_delay,
    &PointToPointNetDevice::Receive, m_link[wire].m_dst, p->Copy());

  // Call the tx anim callback on the net device
  m_txrxPointToPoint(p, src, ...)
  return true;
}
```

an NS-3 event handler

```
bool PointToPointChannel::TransmitStart(
    Ptr<const Packet> p,
    Ptr<PointToPointNetDevice> src,
    Time txTime
) {
  // ...
  uint32_t wire = src == m_link[0].m_src ? 0 : 1;

  Simulator::ScheduleWithContext(
    m_link[wire].m_dst->GetNode()->GetId(),
    txTime + m_delay,
    &PointToPointNetDevice::Receive, m_link[wire].m_dst, p->Copy());

  // Call the tx anim callback on the net device
  m_txrxPointToPoint(p, src, ...)
  return true;
}
```

setup future event; args:
context (for logging mostly)
time event will trigger
method to run + arguments to pass

sample NS-3 simulation — setup (1)

ns3/examples/tcp/tcp-bulk-send.cc

```
// nodes = routers or endpoints
NodeContainer nodes;
nodes.Create(2);

// create simulated point-to-point link
    // also supported: multi-access links
PointToPointHelper pointToPoint;
pointToPoint.SetDeviceAttribute("DataRate", StringValue("500Kbps"));
pointToPoint.SetChannelAttribute("Delay", StringValue("5ms"));

// setup emulated NICs (which have queues, etc.)
NetDeviceContainer devices;
devices = pointToPoint.Install(nodes);

// setup emulated TCP/IP implementation
InternetStackHelper internet;
internet.Install(nodes);

...
```

sample NS-3 simulation — setup (2)

ns3/examples/tcp/tcp-bulk-send.cc

```
// Simulated "applications" that send/receive data
BulkSendHelper source("ns3::TcpSocketFactory", InetSocketAddress(i.GetAddress(1),
source.SetAttribute("MaxBytes", UintegerValue(10000000));
ApplicationContainer sourceApps = source.Install(nodes.Get(0));
sourceApps.Start(Seconds(0));
sourceApps.Stop(Seconds(10));
PacketSinkHelper sink("ns3::TcpSocketFactory", InetSocketAddress(Ipv4Address::GetA
ApplicationContainer sinkApps = sink.Install(nodes.Get(1));
sinkApps.Start(Seconds(0));
sinkApps.Stop(Seconds(10));
```

sample NS-3 simulation — setup (3)

ns3/examples/tcp/tcp-bulk-send.cc

```
AsciiTraceHelper ascii;

// produces text trace file of simulator events
pointToPoint.EnableAsciiAll(ascii.CreateFileStream("tcp-bulk-send.tr"));

// produces PCAP files you can open in Wireshark
pointToPoint.EnablePcapAll("tcp-bulk-send", false);
```

running example: –help

```
$ ./ns3 run examples/tcp/tcp-bulk-send -- --help
[0/2] Re-checking globbed directories...
ninja: no work to do.
tcp-bulk-send [Program Options] [General Arguments]

Program Options:
    --tracing:   Flag to enable/disable tracing [false]
    --maxBytes:  Total number of bytes for application to send [0]

General Arguments:
    --PrintGlobals:              Print the list of globals.
    --PrintGroups:               Print the list of groups.
    --PrintGroup=[group]:        Print all TypeIds of group.
    --PrintTypeIds:              Print all TypeIds.
    --PrintAttributes=[typeid]:  Print all attributes of typeid.
    --PrintVersion:              Print the ns-3 version.
    --PrintHelp:                 Print this help message.
```

running example

```
$ ./ns3 run examples/tcp/tcp-bulk-send -- --maxBytes=1000000 --tracing=true
[0/2] Re-checking globbed directories...
ninja: no work to do.
Total Bytes Received: 566016

$ wireshark tcp-bulk-send-0-0.pcap
$ wireshark tcp-bulk-send-1-0.pcap
$ wireshark tcp-bulk-send-2-0.pcap
$ wireshark tcp-bulk-send-2-1.pcap
```

---

packet traces for each simulated interface: (host)-(interface).pcap

also .tr file with all events

only 566016 bytes received because 10 second simulation only

500Kbit/s = 62KByte/s @ 10 s = about 620KByte

TCPSocketState

class to track socket state

- regarldess of congestion control algorithm

includes (notably for us)

- ‘congesiton state’
- congestion window (cwnd)
- slow start threshold (ssthresh)

TCP congestion states

OPEN — normal

DISORDER — dupacks/SACKs below threshold for recovery

RECOVERY — retransmitting due to dup-ACK/simple SACK

temporarily inflated window to account for dropped packets

LOSS — retransmitting due to timeout/etc.

configurable TCP behavior

TCPRecoveryOps: handle recovery

- EnterRecovery (called on start of retransmitting)
  - default: cwnd $\rightarrow$ ssthresh
- DoRecovery (called for ACK when retransmitting)
  - default: cwnd $\rightarrow$ cwnd + 1 packet
- ExitRecovery (called when 'back to normal')
  - default: cwnd $\rightarrow$ ssthresh

TCPCongestionOps:

- IncreaseWindow (called when new segments ACKed)
- GetSsThresh (called after loss)

assingment TCP changes

will customize a `MyTcpCongestionOps`

...which inherits from TCPNewReno

you can customize increase (IncreaseWindow) and decrease (GetSsThresh)

IncreaseWindow

```
void
TcpNewReno::IncreaseWindow(Ptr<TcpSocketState> tcb, uint32_t segmentsAcked)
{
    NS_LOG_FUNCTION(this << tcb << segmentsAcked);

    if (tcb->m_cWnd < tcb->m_ssThresh)
    {
        segmentsAcked = SlowStart(tcb, segmentsAcked);
    }

    if (tcb->m_cWnd >= tcb->m_ssThresh)
    {
        CongestionAvoidance(tcb, segmentsAcked);
    }
}
```

CongestionAvoidance

```
void
TcpNewReno::CongestionAvoidance(Ptr<TcpSocketState> tcb, uint32_t segmentsAcked)
{
    NS_LOG_FUNCTION(this << tcb << segmentsAcked);

    if (segmentsAcked > 0)
    {
        double adder =
            static_cast<double>(tcb->m_segmentSize * tcb->m_segmentSize) / tcb->m_
        adder = std::max(1.0, adder);
        tcb->m_cWnd += static_cast<uint32_t>(adder);
        ...
    }
}
```

additive increase (plus 1 segment per cwnd)

SlowStart

```
uint32_t
TcpNewReno::SlowStart(Ptr<TcpSocketState> tcb, uint32_t segmentsAcked)
{
    NS_LOG_FUNCTION(this << tcb << segmentsAcked);

    if (segmentsAcked >= 1)
    {
        tcb->m_cWnd += tcb->m_segmentSize;
        ...
        return segmentsAcked - 1;
    }

    return 0;
}
```

multiplicative increase — doubling

GetSsThresh

```
uint32_t
TcpNewReno::GetSsThresh(Ptr<const TcpSocketState> state, uint32_t bytesInFlight)
{
    NS_LOG_FUNCTION(this << state << bytesInFlight);

    return std::max(2 * state->m_segmentSize, bytesInFlight / 2);
}
```

multiplicative decrease (bytesInFlight / 2)

aside: more advanced congestion hooks

TcpCongestionOps also has access to

- ECN (explicit congestion notification) info
- timestamps
- exactly when duplicate ACKs below threshold/ recovery/etc. occurs

used by more advanced TCP implementatoins

assignment preview

start with `MyTcpCongestionOps`

- basically copy of NewRenoTcp

setup to take 'mode' parameter

- mode 0: default behavior
- mode 1 (added by you): take more bandwidth
- mode 2 (added by you): take less bandwidth

NS-3 logging

NS3 logging tools

`https://www.nsnam.org/docs/manual/html/logging-asserts.html`

doing logging:

```
// at top of file
NS_LOG_COMPONENT_DEFINE("SomeLabel");
...
// when you want to log
NS_LOG_INFO("value of x=" << x);
    // INFO could be DEBUG, ERROR, ... too
```

enabling logging:

in C++: `LogComponentEnable("SomeLabel", LOG_LEVEL_DEBUG)`

backup slides